

PROFILING, OPTIMIZATION, PARALLELIZATION

CHRISTOPH PAULIK

OPTIMIZATION

*We should forget about small efficiencies,
say about 97% of the time: premature
optimization is the root of all evil.*

– Donald Knuth

The problem with premature optimization is that you never know in advance where the bottlenecks will be.

WHEN TO OPTIMIZE

- Make it work.
- Make it right
- Make everything work.
- Make everything right.
- Profile Before Optimizing.
- Make it fast. You maintained unit tests, right? Then you can refactor the code mercilessly in order to improve the performance.

Guillermo Schwarz

WHAT IS GENERALLY SLOW IN PYTHON

- CPU bound tasks
- e.g. Nested for loops over large arrays
- But this does not mean that this is the reason why your program is slow
- Can be quite surprising at times

AMDAHL'S LAW

A parallelized program can only be as fast as the slowest single threaded piece of code.

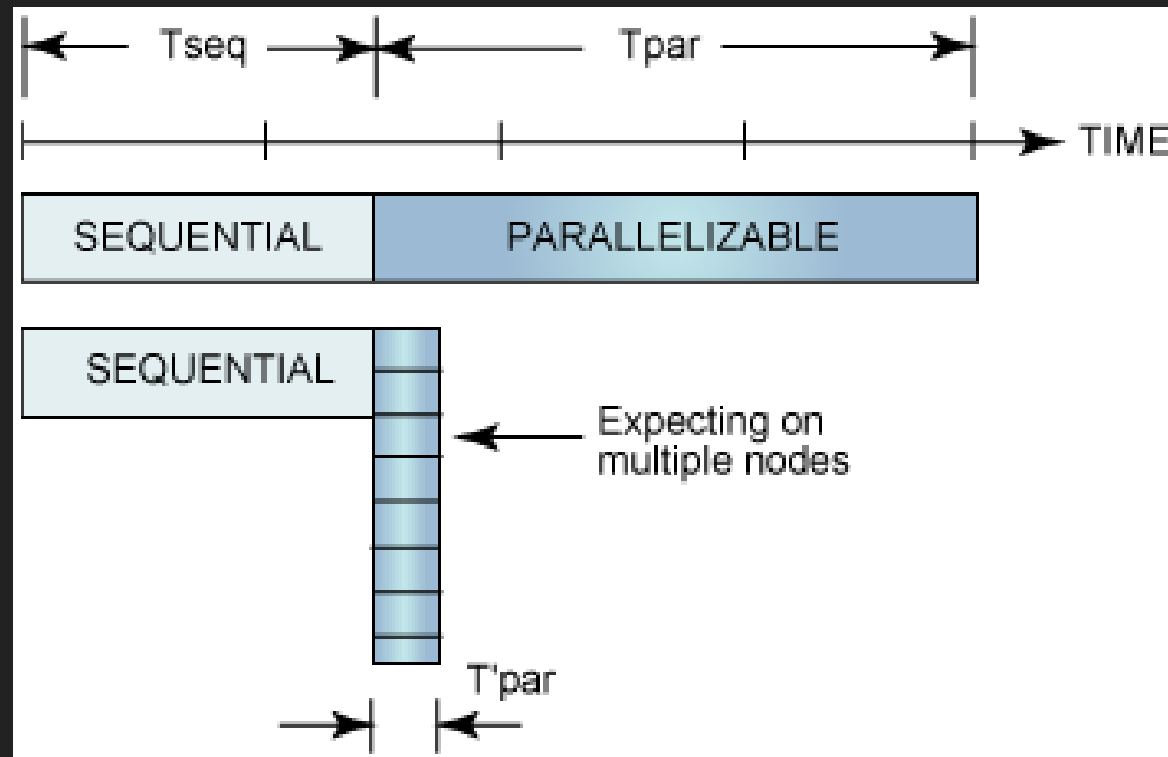


Figure 1: Illustration of Amdahl's law from [IBM](#)

PROFILING

- Built in Python profiler can be called from command line or from Python itself or inside a IPython Notebook.

```
python -m cProfile -o <output> <script-name> <options>
```

- Viewing profile

```
python -m pstats <output>  
runsake <output>
```

[Runsakerun](#) requires wxpython – not portable

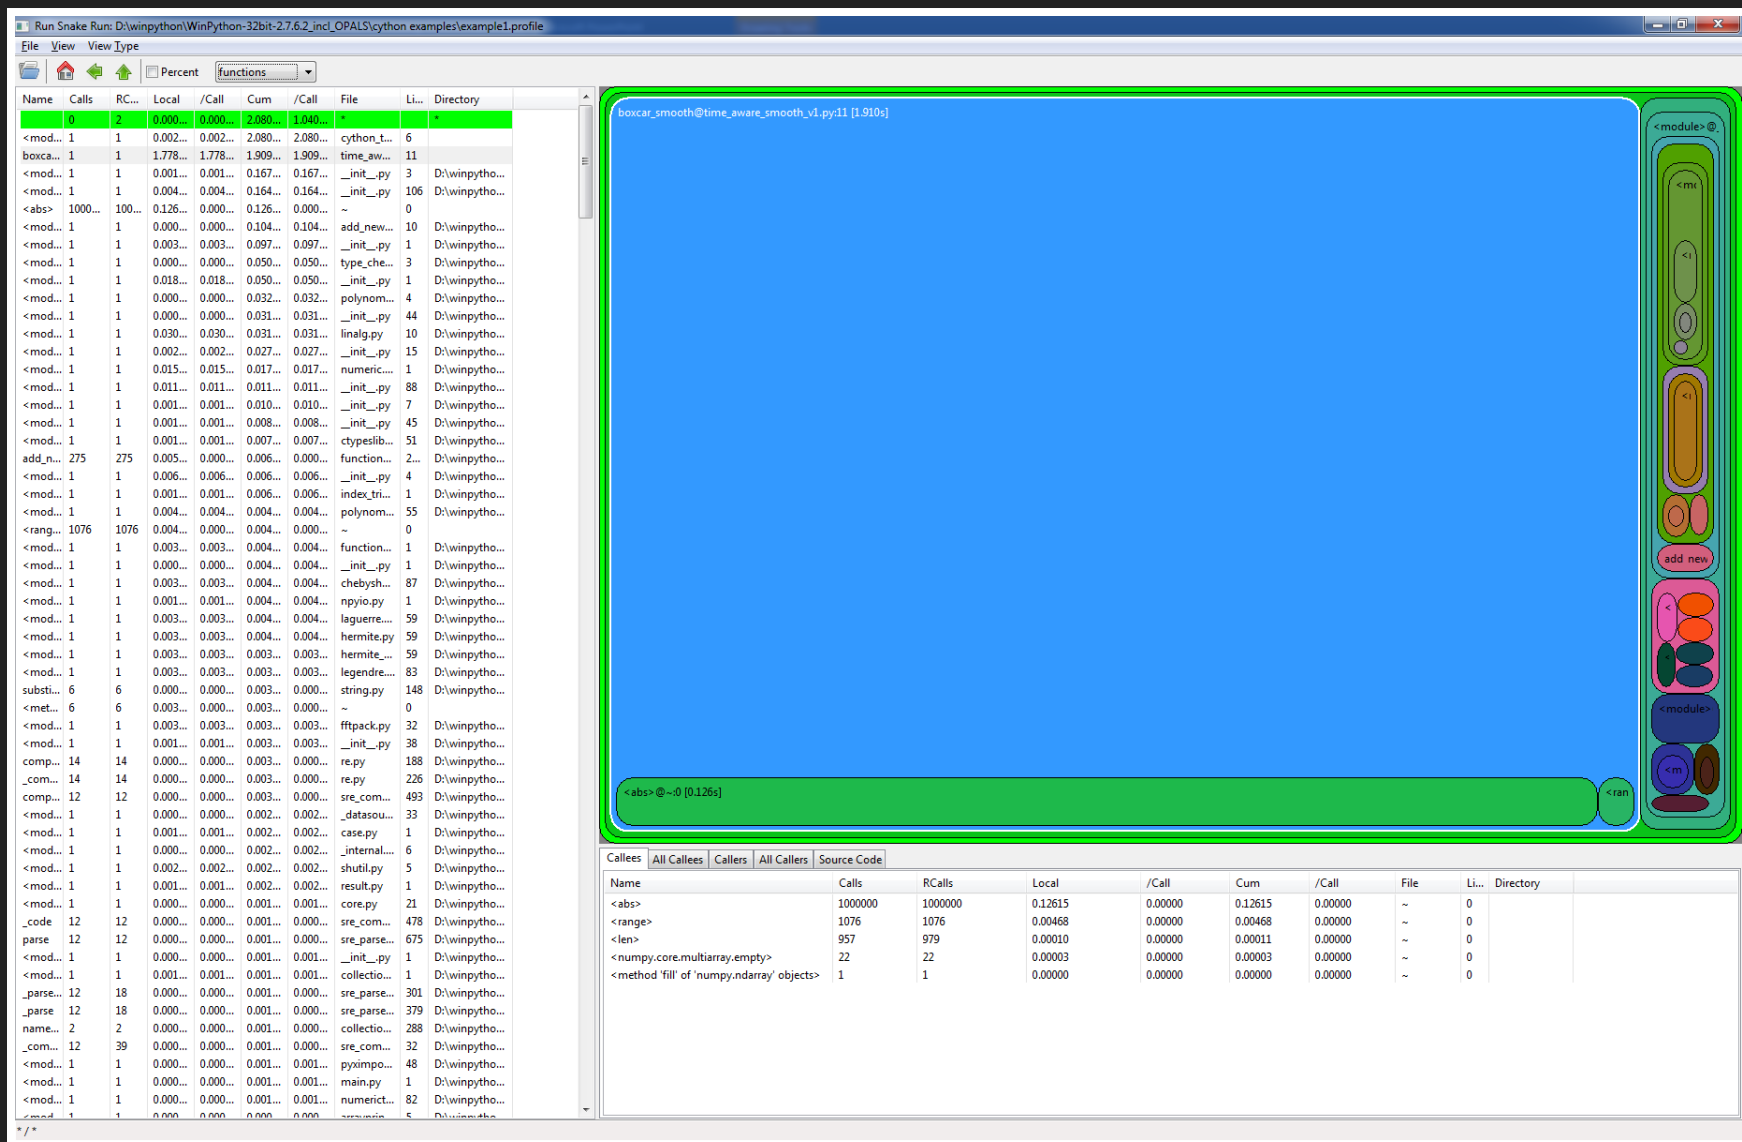


Figure 2: Runsnake interface

LINE PROFILER

- `@profile` decorator for functions that we want to look at

```
kernprof -l -o <outputfilename> <script-name> <options>  
python -m line_profiler <outputfilename>
```


MORE INFORMATION ABOUT PROFILING

- [Python Profilers](#)
- [Line Profiler](#)
- [Tutorial](#)
- [Pstats](#)

IMPROVING PERFORMANCE

Now that you have found the slow part, what to do?

- Can it be done in vectorized form? – numpy? – **no!**
- Are there other packages or existing libraries that do it? – **no!**
- Did you do it in an inefficient way? – **no!**
- Is it CPU bound and not I/O or memory bound? – **yes!**

If those are your answers:

- **Cython** is a good way to go if you do not know C or C++.
- **Numba** should be able to do it automatically. (use Anaconda)

CYTHON

- Cython gives you the combined power of Python and C to let you
 - write Python code that calls back and forth from and to C or C++ code natively at any point.
 - easily tune readable Python code into plain C performance by adding static type declarations. – different ways, we'll focus on one.
 - [Windows install instructions](#)

EXAMPLE

```
def f(x):  
    return x ** 2 - x  
  
def integrate_f(a, b, N):  
    s = 0  
    dx = (b - a) / N  
    for i in range(N):  
        s += f(a + i * dx)  
    return s * dx
```

CYTHON

```
def f(double x):  
    return x ** 2 - x  
  
def integrate_f(double a, double b, int N):  
    cdef int i  
    cdef double s, dx  
    s = 0  
    dx = (b - a / N)  
    for i in range(N):  
        s += f(a + i * dx)  
    return s * dx
```

BUILDING CYTHON CODE

- save it in a `.pyx` file
- use `pyximport` instead of regular `import`
- run the `cython` command line program and then compile the `.c` file manually
- use IPython Notebook

More details in the documentation

EXAMPLE

- open `profiling.ipynb`
- profiling
- find slow part
- make it faster using cython
- IPython parallelization